

Artificial Neural Network (ANN) for predicting Heart Attack

[Code ▾](#)[Hide](#)

```
# Importing the dataset
dataset = read.csv('Heart.csv')
dataset = data.frame(dataset[-1])
dim(dataset)
```

```
[1] 303  14
```

[Hide](#)

```
str(dataset)
```

```
'data.frame':  303 obs. of  14 variables:
 $ Age      : int  63 67 67 37 41 56 62 57 63 53 ...
 $ Sex      : int   1 1 1 1 0 1 0 0 1 1 ...
 $ ChestPain: Factor w/ 4 levels "asymptomatic",...: 4 1 1 2 3 3 1 1 1 1 ...
 $ RestBP   : int  145 160 120 130 130 120 140 120 130 140 ...
 $ Chol     : int  233 286 229 250 204 236 268 354 254 203 ...
 $ Fbs      : int   1 0 0 0 0 0 0 0 0 1 ...
 $ RestECG  : int   2 2 2 0 2 0 2 0 2 2 ...
 $ MaxHR    : int  150 108 129 187 172 178 160 163 147 155 ...
 $ ExAng    : int   0 1 1 0 0 0 0 1 0 1 ...
 $ Oldpeak  : num   2.3 1.5 2.6 3.5 1.4 0.8 3.6 0.6 1.4 3.1 ...
 $ Slope    : int   3 2 2 3 1 1 3 1 2 3 ...
 $ Ca       : int   0 3 2 0 0 0 2 0 1 0 ...
 $ Thal     : Factor w/ 3 levels "fixed","normal",...: 1 2 3 2 2 2 2 2 3 3 ...
 $ AHD      : Factor w/ 2 levels "No","Yes": 1 2 2 1 1 1 2 1 2 2 ...
```

[Hide](#)

```
# Checking for missing data
d3=dataset
for(i in 1:ncol(d3))
{
  print(colnames(d3[i]))
  print(sum(is.na(d3[i])))
}
```

```
[1] "Age"  
[1] 0  
[1] "Sex"  
[1] 0  
[1] "ChestPain"  
[1] 0  
[1] "RestBP"  
[1] 0  
[1] "Chol"  
[1] 0  
[1] "Fbs"  
[1] 0  
[1] "RestECG"  
[1] 0  
[1] "MaxHR"  
[1] 0  
[1] "ExAng"  
[1] 0  
[1] "Oldpeak"  
[1] 0  
[1] "Slope"  
[1] 0  
[1] "Ca"  
[1] 4  
[1] "Thal"  
[1] 2  
[1] "AHD"  
[1] 0
```

Hide

```
# Some missing data is present  
# Removing Missing Data from data set  
dataset = na.omit(dataset)  
dim(dataset)
```

```
[1] 297  14
```

Hide

```
# Checking for missing data again  
d3=dataset  
for(i in 1:ncol(d3))  
{  
  print(colnames(d3[i]))  
  print(sum(is.na(d3[i])))  
}
```

```
[1] "Age"  
[1] 0  
[1] "Sex"  
[1] 0  
[1] "ChestPain"  
[1] 0  
[1] "RestBP"  
[1] 0  
[1] "Chol"  
[1] 0  
[1] "Fbs"  
[1] 0  
[1] "RestECG"  
[1] 0  
[1] "MaxHR"  
[1] 0  
[1] "ExAng"  
[1] 0  
[1] "Oldpeak"  
[1] 0  
[1] "Slope"  
[1] 0  
[1] "Ca"  
[1] 0  
[1] "Thal"  
[1] 0  
[1] "AHD"  
[1] 0
```

Hide

```
# Encoding the categorical & numerical variables in dataset  
dataset$Age = as.numeric(dataset$Age)  
dataset$Sex = as.factor(dataset$Sex)  
dataset$ChestPain = as.factor(dataset$ChestPain)  
dataset$RestBP = as.numeric(dataset$RestBP)  
dataset$Chol = as.numeric(dataset$Chol)  
dataset$Fbs = as.factor(dataset$Fbs)  
dataset$RestECG = as.factor(dataset$RestECG)  
dataset$MaxHR = as.numeric(dataset$MaxHR)  
dataset$ExAng = as.factor(dataset$ExAng)  
dataset$Oldpeak = as.numeric(dataset$Oldpeak)  
dataset$Slope = as.factor(dataset$Slope)  
dataset$Ca = as.factor(dataset$Ca)  
dataset$Thal = as.factor(dataset$Thal)  
dataset$AHD = ifelse(dataset$AHD == "Yes",1,0)  
dataset$AHD = as.factor(dataset$AHD)  
str(dataset)
```

```
'data.frame': 297 obs. of 14 variables:
 $ Age      : num  63 67 67 37 41 56 62 57 63 53 ...
 $ Sex      : Factor w/ 2 levels "0","1": 2 2 2 2 1 2 1 1 2 2 ...
 $ ChestPain: Factor w/ 4 levels "asymptomatic",...: 4 1 1 2 3 3 1 1 1 1 ...
 $ RestBP   : num  145 160 120 130 130 120 140 120 130 140 ...
 $ Chol     : num  233 286 229 250 204 236 268 354 254 203 ...
 $ Fbs      : Factor w/ 2 levels "0","1": 2 1 1 1 1 1 1 1 1 2 ...
 $ RestECG  : Factor w/ 3 levels "0","1","2": 3 3 3 1 3 1 3 1 3 3 ...
 $ MaxHR    : num  150 108 129 187 172 178 160 163 147 155 ...
 $ ExAng    : Factor w/ 2 levels "0","1": 1 2 2 1 1 1 1 2 1 2 ...
 $ Oldpeak  : num  2.3 1.5 2.6 3.5 1.4 0.8 3.6 0.6 1.4 3.1 ...
 $ Slope    : Factor w/ 3 levels "1","2","3": 3 2 2 3 1 1 3 1 2 3 ...
 $ Ca       : Factor w/ 4 levels "0","1","2","3": 1 4 3 1 1 1 3 1 2 1 ...
 $ Thal     : Factor w/ 3 levels "fixed","normal",...: 1 2 3 2 2 2 2 2 3 3 ...
 $ AHD      : Factor w/ 2 levels "0","1": 1 2 2 1 1 1 2 1 2 2 ...
 - attr(*, "na.action")=Class 'omit' Named int [1:6] 88 167 193 267 288 303
 .. ...- attr(*, "names")= chr [1:6] "88" "167" "193" "267" ...
```

Hide

```
# Splitting the dataset into the Training set and Test set
library(caTools)
set.seed(123)
split = sample.split(dataset$AHD, SplitRatio = 0.75)
training_set = subset(dataset, split == TRUE)
test_set = subset(dataset, split == FALSE)
# Feature Scaling
training_set[,c(1,4,5,8,10)] = scale(training_set[,c(1,4,5,8,10)])
test_set[,c(1,4,5,8,10)] = scale(test_set[,c(1,4,5,8,10)])
```

Hide

```
# Fitting ANN to the Training set
# Nodes in hidden layers selected by Avg (# Input + # Output Nodes) = (14+1)/2 = 7.5; 8
#install.packages('h2o')
library(h2o)
h2o.init(ip = "localhost")
```

Connection successful!

R is connected to the H2O cluster:

```
H2O cluster uptime:      23 minutes 24 milliseconds
H2O cluster timezone:    America/New_York
H2O data parsing timezone: UTC
H2O cluster version:     3.20.0.2
H2O cluster version age:  1 month and 12 days
H2O cluster name:        H2O_started_from_R_Rudrendu_xai688
H2O cluster total nodes: 1
H2O cluster total memory: 1.73 GB
H2O cluster total cores: 4
H2O cluster allowed cores: 4
H2O cluster healthy:     TRUE
H2O Connection ip:       localhost
H2O Connection port:     54321
H2O Connection proxy:    NA
H2O Internal Security:   FALSE
H2O API Extensions:      Algos, AutoML, Core V3, Core V4
R Version:               R version 3.4.4 (2018-03-15)
```

Hide

```
model = h2o.deeplearning(y = 14,
                          training_frame = as.h2o(training_set),
                          activation = 'Rectifier',
                          hidden = c(8,8),
                          epochs = 100,
                          train_samples_per_iteration = -2)
```

```
|
|
|
|=====| 0%
|=====| 100%

|
|
|
|=====| 0%
|=====| 100%
```

Hide

```
# Predicting the Test set results
y_pred = h2o.predict(model, newdata = as.h2o(test_set[-14]))
```

```

|
|
|
|=====| 0%
|=====| 100%

|
|
|
|=====| 0%
|=====| 100%

```

Hide

```
str(y_pred)
```

```

Class 'H2OFrame' <environment: 0x000000002a96578>
- attr(*, "op")= chr "predictions_b362_DeepLearning_model_R_1532814754538_14_on_file31e475b5a68_sid_b064_15"
- attr(*, "id")= chr "predictions_b362_DeepLearning_model_R_1532814754538_14_on_file31e475b5a68_sid_b064_15"
- attr(*, "eval")= logi FALSE
- attr(*, "nrow")= int 74
- attr(*, "ncol")= int 3
- attr(*, "types")=List of 3
..$ : chr "enum"
..$ : chr "real"
..$ : chr "real"
- attr(*, "data")='data.frame': 10 obs. of 3 variables:
..$ predict: Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 2
..$ p0 : num 0.999 0.997 0.96 0.908 0.982 ...
..$ p1 : num 0.000569 0.003035 0.03998 0.092209 0.017695 ...

```

Hide

```
dim(test_set)
```

```
[1] 74 14
```

Hide

```

y_pred=y_pred[1]
y_pred = as.vector(y_pred)
# Making the Confusion Matrix
cm = table(test_set[, 14], y_pred)
print(cm)

```

```

y_pred
  0  1
0 34  6
1  8 26

```

Hide

```
Model_Accuracy=(cm[1,1]+cm[2,2])/(cm[1,1]+cm[1,2]+cm[2,1]+cm[2,2])  
print(" Model Accuracy is")
```

```
[1] " Model Accuracy is"
```

Hide

```
print(Model_Accuracy)
```

```
[1] 0.8108108
```

Hide

```
# Model Accuracy is 81% which seems quite good.  
h2o.shutdown()
```

...